

# **Monte Carlo Proxy Pricing**

A textbook-style methodology for European, American, Asian, barrier, cliquet, basket Asian, SLV cliquet, basket cliquet, autocallable, and random-payoff proxies

Proxy Pricing Research Project

July 11, 2026

# Contents

|                                                     |           |
|-----------------------------------------------------|-----------|
| <b>Notation and Conventions</b>                     | <b>vi</b> |
| <b>1 The Pricing Problem</b>                        | <b>1</b>  |
| 1.1 Risk-neutral value . . . . .                    | 1         |
| 1.2 What makes a good proxy state . . . . .         | 1         |
| 1.3 Error decomposition . . . . .                   | 2         |
| <b>2 Monte Carlo Label Construction</b>             | <b>3</b>  |
| 2.1 Plain estimator . . . . .                       | 3         |
| 2.2 Sobol and antithetics . . . . .                 | 3         |
| 2.3 Default simulation policy . . . . .             | 4         |
| 2.4 From Sobol uniforms to GBM paths . . . . .      | 4         |
| 2.5 Likelihood-ratio importance sampling . . . . .  | 4         |
| 2.6 Mixture importance sampling . . . . .           | 5         |
| 2.7 Control variates . . . . .                      | 5         |
| <b>3 Proxy Fitting Principles</b>                   | <b>6</b>  |
| 3.1 Targets . . . . .                               | 6         |
| 3.2 Exact regions . . . . .                         | 6         |
| 3.3 Relative error floor . . . . .                  | 6         |
| <b>4 One-Dimensional Shape-Preserving Proxies</b>   | <b>7</b>  |
| 4.1 Why PCHIP is the default for 1D . . . . .       | 7         |
| 4.2 Cubic Hermite form . . . . .                    | 7         |
| 4.3 PCHIP slope construction . . . . .              | 7         |
| 4.4 Why exact interpolation is acceptable . . . . . | 8         |
| 4.5 Akima, splines, and Bernstein . . . . .         | 8         |

|          |                                                |           |
|----------|------------------------------------------------|-----------|
| <b>5</b> | <b>Sparse Chebyshev and Higher Dimension</b>   | <b>9</b>  |
| 5.1      | Chebyshev scaling . . . . .                    | 9         |
| 5.2      | Ridge regression . . . . .                     | 9         |
| 5.3      | Sparse multi-index basis . . . . .             | 9         |
| 5.4      | Baseline plus residual . . . . .               | 10        |
| 5.5      | Why sparsity is necessary . . . . .            | 10        |
| <b>6</b> | <b>Instrument Recipes</b>                      | <b>11</b> |
| 6.1      | Recipe template . . . . .                      | 11        |
| 6.2      | European option . . . . .                      | 11        |
| 6.3      | Arithmetic Asian option . . . . .              | 12        |
| 6.4      | Ten-asset basket Asian . . . . .               | 12        |
| 6.5      | American and Bermudan options . . . . .        | 12        |
| 6.6      | Barrier option . . . . .                       | 13        |
| 6.7      | Cliquet and SLV cliquet . . . . .              | 13        |
| 6.8      | Autocallable and random payoff . . . . .       | 13        |
| <b>7</b> | <b>Contract Definitions and Payoff Algebra</b> | <b>14</b> |
| 7.1      | European call and put . . . . .                | 14        |
| 7.2      | Arithmetic Asian call . . . . .                | 14        |
| 7.3      | Basket Asian call . . . . .                    | 15        |
| 7.4      | Barrier call . . . . .                         | 15        |
| 7.5      | Cliquet payoff . . . . .                       | 15        |
| 7.6      | Basket cliquet coupon variants . . . . .       | 16        |
| 7.7      | Autocallable payoff recursion . . . . .        | 16        |
| <b>8</b> | <b>Implementation Algorithms</b>               | <b>17</b> |
| 8.1      | Generic proxy construction algorithm . . . . . | 17        |
| 8.2      | European algorithm . . . . .                   | 17        |
| 8.3      | Asian algorithm . . . . .                      | 18        |
| 8.4      | Basket Asian algorithm . . . . .               | 18        |
| 8.5      | American and Bermudan algorithm . . . . .      | 18        |
| 8.6      | Barrier algorithm . . . . .                    | 19        |
| 8.7      | Cliquet and basket cliquet algorithm . . . . . | 19        |

|          |                                                                  |           |
|----------|------------------------------------------------------------------|-----------|
| <b>9</b> | <b>Validation and Timing</b>                                     | <b>20</b> |
| 9.1      | Independent benchmark principle . . . . .                        | 20        |
| 9.2      | Current reporting . . . . .                                      | 20        |
| <b>A</b> | <b>Line-by-Line Proofs: Conditional Pricing and State Choice</b> | <b>21</b> |
| A.1      | Risk-neutral pricing formula . . . . .                           | 21        |
| A.2      | When a proxy state is sufficient . . . . .                       | 21        |
| A.3      | Tower property for backward induction . . . . .                  | 22        |
| <b>B</b> | <b>Line-by-Line Proofs: Monte Carlo</b>                          | <b>23</b> |
| B.1      | Unbiased sample mean . . . . .                                   | 23        |
| B.2      | Likelihood-ratio shift . . . . .                                 | 23        |
| B.3      | Mixture likelihood ratio . . . . .                               | 24        |
| B.4      | Antithetic estimator . . . . .                                   | 24        |
| <b>C</b> | <b>Line-by-Line Proofs: GBM and European Options</b>             | <b>25</b> |
| C.1      | GBM solution . . . . .                                           | 25        |
| C.2      | Black-Scholes call . . . . .                                     | 25        |
| <b>D</b> | <b>Line-by-Line Proofs: Asian State Reduction</b>                | <b>26</b> |
| D.1      | Adjusted moneyness . . . . .                                     | 26        |
| D.2      | Geometric Asian control variate . . . . .                        | 27        |
| <b>E</b> | <b>Line-by-Line Proofs: Barriers</b>                             | <b>28</b> |
| E.1      | In/out parity . . . . .                                          | 28        |
| E.2      | Brownian-bridge lower-barrier survival . . . . .                 | 28        |
| <b>F</b> | <b>Line-by-Line Proofs: Exercise and Cliquet Bounds</b>          | <b>29</b> |
| F.1      | American recursion . . . . .                                     | 29        |
| F.2      | Cliquet exact tails . . . . .                                    | 29        |
| <b>G</b> | <b>Line-by-Line Proofs: PCHIP</b>                                | <b>30</b> |
| G.1      | Unique Hermite cubic . . . . .                                   | 30        |
| G.2      | Why PCHIP avoids artificial extrema . . . . .                    | 30        |
| <b>H</b> | <b>Line-by-Line Proofs: Sparse Chebyshev Ridge</b>               | <b>32</b> |

|          |                                                           |           |
|----------|-----------------------------------------------------------|-----------|
| H.1      | Chebyshev boundedness . . . . .                           | 32        |
| H.2      | Ridge normal equation . . . . .                           | 32        |
| <b>I</b> | <b>Line-by-Line Proofs: Bounded Targets and Metrics</b>   | <b>33</b> |
| I.1      | Logit bounded target . . . . .                            | 33        |
| I.2      | Relative-error floor . . . . .                            | 33        |
| <b>J</b> | <b>Line-by-Line Proofs: Basket Moments and PCA</b>        | <b>34</b> |
| J.1      | Basket Asian first moment . . . . .                       | 34        |
| J.2      | Basket Asian second moment . . . . .                      | 34        |
| J.3      | Moment-matched lognormal baseline . . . . .               | 35        |
| J.4      | PCA optimality . . . . .                                  | 35        |
| <b>K</b> | <b>Line-by-Line Proofs: Sobol, Design, and Validation</b> | <b>36</b> |
| K.1      | Why powers of two are used . . . . .                      | 36        |
| K.2      | Common random numbers . . . . .                           | 36        |
| K.3      | Benchmark independence . . . . .                          | 36        |
| <b>L</b> | <b>Line-by-Line Proofs: SLV and Basket Cliquet</b>        | <b>38</b> |
| L.1      | Variance process discretization . . . . .                 | 38        |
| L.2      | Local volatility multiplier . . . . .                     | 38        |
| L.3      | Grouped accrued labels . . . . .                          | 38        |
| L.4      | Why basket order statistics are hard . . . . .            | 39        |
|          | <b>References</b>                                         | <b>40</b> |

# Reader's Guide

This document is written as a second-year undergraduate text. It assumes calculus, linear algebra, basic probability, and a willingness to read a few stochastic-process formulas. The goal is not to hide the implementation behind jargon. The goal is to show exactly how a Monte Carlo price label becomes a fast proxy.

The main chapters explain the workflow. The appendices give line-by-line proofs for the formulas used in the code. Every nontrivial implementation trick is classified as one of:

- an exact pricing identity,
- an unbiased Monte Carlo variance-reduction identity,
- a numerical approximation whose error must be validated.

# Notation and Conventions

The same notation is used throughout the document.

| Symbol          | Meaning                                                                                                               |
|-----------------|-----------------------------------------------------------------------------------------------------------------------|
| $S_t$           | Spot price at time $t$ . For baskets, $S_t^{(j)}$ is asset $j$ .                                                      |
| $r, q, \sigma$  | Risk-free rate, dividend yield, and volatility in GBM examples.                                                       |
| $\tau$          | Remaining time to maturity, usually $T - t$ .                                                                         |
| $H$             | Final payoff or discounted cash-flow functional before discounting.                                                   |
| $V(t, x)$       | Model value at time $t$ when the pricing state is $x$ .                                                               |
| $X_t$           | State vector used by the proxy.                                                                                       |
| $Z$             | Standard-normal simulation shocks after inverse-CDF mapping from Sobol points.                                        |
| $N$             | Number of Monte Carlo paths or number of fixings, depending on context. The surrounding formula identifies which one. |
| $A_i, R_i, a_i$ | Running sums: Asian running sum, basket running sum, and cliquet accrued return.                                      |
| $\ell, u$       | Local floor and cap for a cliquet coupon.                                                                             |
| $G_L, G_U$      | Global floor and cap for a cliquet payoff.                                                                            |
| $\eta$          | Relative-error denominator floor.                                                                                     |

All prices in the experiments are model prices. If the simulator is GBM, SLV, or a tree, the proxy is judged against that model. A proxy can be accurate relative to the model and the model can still be a poor market model. The two questions are deliberately separated.

# Chapter 1

## The Pricing Problem

### 1.1 Risk-neutral value

Let  $H$  be the payoff of a derivative at maturity or final settlement time  $T$ . Under the risk-neutral measure  $\mathbb{Q}$ , the time- $t$  value is

$$V_t = B_t \mathbb{E}_{\mathbb{Q}} \left[ \frac{H}{B_T} \middle| \mathcal{F}_t \right],$$

where  $B_t = \exp(\int_0^t r_s ds)$  is the money-market account. If the short rate is constant,

$$V_t = e^{-r(T-t)} \mathbb{E}_{\mathbb{Q}}[H \mid \mathcal{F}_t].$$

In the code, the expensive object is the conditional expectation. The proxy is a fast approximation of the map

$$x \longmapsto V(t, x),$$

where  $x$  is the state available at the valuation date.

### 1.2 What makes a good proxy state

A state is good when it is sufficient, low-dimensional, and scaled sensibly.

**Definition 1.1** (Sufficient state). A vector  $X_t$  is sufficient for pricing if, after conditioning on  $X_t$ , the distribution of all future cash flows no longer depends on earlier path history.

Examples:

- European option:  $X_t = S_t$ .
- Arithmetic Asian option:  $X_t = (S_t, A_t)$ , where  $A_t$  is the running sum of past fixings.
- American put:  $X_t = S_t$ , but the time index matters because of remaining exercise dates.
- Barrier option:  $X_t = (S_t, I_t)$ , where  $I_t$  says whether the barrier has already been hit.
- Single-name SLV cliquet:  $X_t = (a_t, S_t, v_t)$ , accrued coupon, spot, and variance.
- Basket SLV cliquet:  $X_t = (a_t, S_t^{(1:3)}, v_t^{(1:3)})$ , plus contract variant.



### 1.3 Error decomposition

The observed proxy error is not one number with one cause. A useful decomposition is

$$\widehat{V}_{\text{proxy}} - V_{\text{true}} = \underbrace{\widehat{V}_{\text{proxy}} - V_{\text{MC-label}}}_{\text{fitting error}} + \underbrace{V_{\text{MC-label}} - V_{\text{model}}}_{\text{training MC noise}} + \underbrace{V_{\text{model}} - V_{\text{true}}}_{\text{model/discretization error}}.$$

The experiments in this repository mostly study the first two terms under fixed model assumptions.

## Chapter 2

# Monte Carlo Label Construction

### 2.1 Plain estimator

For a state  $x$ , define the discounted path value

$$Y(x, Z) = e^{-r\tau} H(x, Z),$$

where  $Z$  represents all future random shocks. The target label is

$$V(x) = \mathbb{E}[Y(x, Z)].$$

With  $N$  simulated paths,

$$\hat{V}_N(x) = \frac{1}{N} \sum_{n=1}^N Y(x, Z_n).$$

The estimator is unbiased when the paths have the correct law. Its standard error is estimated by

$$\widehat{\text{SE}}(\hat{V}_N) = \sqrt{\frac{1}{N(N-1)} \sum_{n=1}^N (Y_n - \bar{Y})^2}.$$

### 2.2 Sobol and antithetics

Sobol points reduce integration error by filling the unit cube more evenly than independent random points. The implementation rounds the number of paths to a power of two because Sobol nets are naturally balanced in blocks of  $2^m$ .

Antithetic pairing uses the fact that, for a standard normal vector  $Z$ ,

$$Z \stackrel{d}{=} -Z.$$

The paired payoff

$$\frac{f(Z) + f(-Z)}{2}$$

has the same expectation as  $f(Z)$ . It often has lower variance for monotone or nearly monotone payoffs.

### 2.3 Default simulation policy

The default policy used by the project is:

1. Use Sobol low-discrepancy points whenever the path dimension is known in advance.
2. Use antithetic pairing when the payoff simulator accepts both  $Z$  and  $-Z$ .
3. Use a likelihood-ratio shift or shift mixture for tail-sensitive labels whenever a useful direction is available.
4. Use the unshifted component as part of the mixture unless a closed-form variance calculation says a pure shift is safe.
5. Use independent randomization or independent Sobol scrambles for benchmarks.

The phrase “whenever possible” matters. For example, if a payoff is synthetic and the useful tail direction is unknown, the safe choice is a mixture containing negative, zero, and positive shifts. If a benchmark uses a finite-difference or tree method, no MC importance sampling is needed because there is no path simulation in that benchmark.

### 2.4 From Sobol uniforms to GBM paths

For a  $d$ -asset GBM model over one step  $\Delta t$ , start with a Sobol vector

$$U = (U_1, \dots, U_d) \in (0, 1)^d.$$

Map it to independent normals by

$$Z_i = \Phi^{-1}(U_i).$$

If  $L$  is a Cholesky factor of the correlation matrix  $\rho$ , so that  $LL^\top = \rho$ , set

$$\tilde{Z} = LZ.$$

Then the simulated next spot is

$$S_{t+\Delta t}^{(j)} = S_t^{(j)} \exp\left((r - q_j - \tfrac{1}{2}\sigma_j^2)\Delta t + \sigma_j\sqrt{\Delta t}\tilde{Z}_j\right).$$

For a path with many dates, repeat this step date by date. The Brownian increments are independent across time before correlation is applied across assets.

### 2.5 Likelihood-ratio importance sampling

State enrichment and importance sampling are different.

State enrichment changes the distribution of training states  $x_i$ . It does not change the path law for any individual conditional expectation.

Importance sampling changes the path density from  $p$  to  $q$  and multiplies by a likelihood ratio:

$$\mathbb{E}_p[h(Z)] = \mathbb{E}_q\left[h(Z)\frac{p(Z)}{q(Z)}\right].$$

For a normal mean shift,  $q(z) = \phi(z - \theta)$ , so

$$\frac{p(z)}{q(z)} = \exp\left(-\theta z + \frac{1}{2}\theta^2\right).$$

This identity is exact; its usefulness depends on variance.

## 2.6 Mixture importance sampling

A single shift can be too opinionated. The more generic tail-biased method is a mixture:

$$q(z) = \sum_{m=1}^M \pi_m \phi_d(z - \theta_m), \quad \pi_m > 0, \quad \sum_{m=1}^M \pi_m = 1.$$

Here  $\phi_d$  is the  $d$ -dimensional standard normal density and  $\theta_m$  is a vector shift. The likelihood ratio is

$$L(z) = \frac{\phi_d(z)}{\sum_{m=1}^M \pi_m \phi_d(z - \theta_m)}.$$

The weighted estimator is

$$\hat{V}_{\text{IS}} = \frac{1}{N} \sum_{n=1}^N h(Z_n) L(Z_n), \quad Z_n \sim q.$$

For option training, mixture shifts are useful because one model can include lower-tail, upper-tail, and central scenarios without needing a different formula for every product. The cost is that likelihood weights can become variable; therefore the code tracks benchmark error and avoids using an aggressive shift as the only source of paths.

## 2.7 Control variates

If  $C$  has known mean  $m_C$ , then

$$Y_\beta = Y - \beta(C - m_C)$$

has the same expectation as  $Y$ . The variance-minimizing coefficient is

$$\beta^\star = \frac{\text{Cov}(Y, C)}{\text{Var}(C)}.$$

Asian options use a geometric Asian control when the arithmetic payoff and geometric payoff are strongly correlated.

## Chapter 3

# Proxy Fitting Principles

### 3.1 Targets

Raw prices can be hard to fit because they range from nearly zero to almost linear. The code uses transformations such as:

$$\log(V + \varepsilon), \quad \log \frac{V + \varepsilon}{B + \varepsilon}, \quad \text{logit} \frac{V - L}{U - L}.$$

Here  $B$  is a baseline, and  $L, U$  are known lower and upper bounds.

### 3.2 Exact regions

Never ask a regressor to learn something known analytically.

Examples:

- At maturity, use the payoff exactly.
- Deep Asian call in-the-money region: the positive part disappears and the value is linear.
- Cliquet global floor/cap tails: the final clipped payoff is already known.
- American/Bermudan exercise: impose  $V = \max(\text{intrinsic}, \text{continuation})$ .

### 3.3 Relative error floor

The reported relative error is

$$e_{\text{rel}} = \frac{|\hat{V} - V|}{\max(|V|, \eta)}.$$

This is ordinary relative error when  $|V| \geq \eta$ , and it avoids meaningless exploding percentages when  $V$  is tiny.

## Chapter 4

# One-Dimensional Shape-Preserving Proxies

### 4.1 Why PCHIP is the default for 1D

For a single feature, PCHIP is a strong operational default:

- it interpolates labels exactly;
- it is local, so one noisy point does not create global oscillation;
- it preserves monotone shape when the data are monotone;
- it behaves well near exercise kinks, barrier transitions, and low-value tails.

### 4.2 Cubic Hermite form

On  $[x_i, x_{i+1}]$ , let  $h_i = x_{i+1} - x_i$ ,  $t = (x - x_i)/h_i$ . A cubic Hermite segment is

$$\begin{aligned} H_i(t) = & y_i(2t^3 - 3t^2 + 1) + h_i d_i(t^3 - 2t^2 + t) \\ & + y_{i+1}(-2t^3 + 3t^2) + h_i d_{i+1}(t^3 - t^2). \end{aligned}$$

PCHIP chooses slopes  $d_i$  from neighboring secants in a way that prevents artificial extrema.

### 4.3 PCHIP slope construction

Let

$$\delta_i = \frac{y_{i+1} - y_i}{x_{i+1} - x_i}$$

be the secant slope. At an interior knot  $x_i$ , PCHIP applies this rule.

- If  $\delta_{i-1}\delta_i \leq 0$ , set  $d_i = 0$ . This prevents a spurious turn at a sign change.

- If  $\delta_{i-1}$  and  $\delta_i$  have the same sign, set

$$d_i = \frac{w_1 + w_2}{w_1/\delta_{i-1} + w_2/\delta_i}, \quad w_1 = 2h_i + h_{i-1}, \quad w_2 = h_i + 2h_{i-1}.$$

The second formula is a weighted harmonic mean. A harmonic mean is pulled toward the smaller magnitude, which stops the interpolant from using a large slope to overshoot a flat part of the curve.

Endpoint slopes use one-sided information and are then clipped to respect the first interior secant. This is why SciPy's `PchipInterpolator` is preferred in the main code: it gives a mature, tested implementation of the Fritsch-Butland endpoint and interior rules.

## 4.4 Why exact interpolation is acceptable

PCHIP interpolates labels exactly, so it is not a smoother. That is acceptable only when MC noise is already controlled by Sobol paths, antithetics, control variates, and enough paths. If labels are visibly noisy, a smoothing spline or local regression can have lower average error. The project therefore uses PCHIP as the default one-feature method but still treats label construction as the first-order accuracy problem.

## 4.5 Akima, splines, and Bernstein

Akima interpolation is also local and robust to outliers. Natural cubic splines are smoother but can overshoot. Bernstein polynomials are shape-friendly on bounded intervals and were competitive for some barrier variants. For one-dimensional production defaults, the current rule is:

PCHIP first, compare Akima/Bernstein when the validation surface says so.

## Chapter 5

# Sparse Chebyshev and Higher Dimension

### 5.1 Chebyshev scaling

Chebyshev polynomials satisfy

$$T_n(\cos \theta) = \cos(n\theta),$$

so  $|T_n(x)| \leq 1$  for  $x \in [-1, 1]$ . This is why features are scaled before fitting:

$$z_j = 2 \frac{x_j - \ell_j}{u_j - \ell_j} - 1.$$

### 5.2 Ridge regression

Given a design matrix  $A$ , labels  $y$ , and penalty matrix  $\Gamma$ , ridge regression solves

$$\min_{\beta} \|A\beta - y\|_2^2 + \lambda \|\Gamma\beta\|_2^2.$$

The normal equation is

$$(A^\top A + \lambda \Gamma^\top \Gamma) \beta = A^\top y.$$

### 5.3 Sparse multi-index basis

For a  $D$ -dimensional state  $z = (z_1, \dots, z_D) \in [-1, 1]^D$ , a tensor Chebyshev basis term is

$$\Phi_\alpha(z) = \prod_{j=1}^D T_{\alpha_j}(z_j), \quad \alpha = (\alpha_1, \dots, \alpha_D).$$

A full tensor basis with  $\alpha_j \leq p$  is usually too large. A total-degree basis keeps only

$$|\alpha|_1 = \alpha_1 + \dots + \alpha_D \leq p.$$



The implementation also allows payoff-aware pruning. For example, a basket Asian correction keeps level, dispersion, running-sum, and a few PCA interaction terms before it considers high-order cross-products. This is not a pricing theorem; it is a bias-variance decision that must be checked by validation.

## 5.4 Baseline plus residual

High-dimensional products are easier to fit when the proxy is not asked to learn the whole price from scratch. The generic decomposition is

$$V(x) = B(x) R(x),$$

or, with an additive residual,

$$V(x) = B(x) + \varepsilon_{\text{res}}(x).$$

Here  $B(x)$  is a cheap baseline: Black-Scholes, moment-matched lognormal, intrinsic value plus discounted floor, or a low-dimensional approximation. The fitted target can then be

$$y(x) = \log \frac{V(x) + \epsilon}{B(x) + \epsilon}.$$

When  $B$  captures most of the shape, the residual is flatter and lower dimensional. This is the main reason the basket Asian and basket cliquet proxies improved after adding moment, PCA, and grouped-label features.

## 5.5 Why sparsity is necessary

A full tensor degree- $p$  basis in  $D$  variables contains  $(p+1)^D$  terms. With  $D = 7$  and  $p = 5$ , that is 279,936 terms before fitting. Basket cliquet therefore needs sparse terms, payoff-aware features, or a slower safety method.

# Chapter 6

## Instrument Recipes

### 6.1 Recipe template

Each product section answers the same six questions:

1. What is the payoff?
2. What state is sufficient for future cash flows?
3. Which coordinate is used by the proxy?
4. How are MC labels generated?
5. What benchmark is used?
6. Which parts are exact identities and which parts are fitted approximations?

This template is important because a proxy can look accurate for the wrong reason. For example, reusing training paths in the benchmark can hide error, and using a coordinate that is not sufficient can look fine on a small grid but fail when parameters change.

### 6.2 European option

State:  $S$ . Coordinate:

$$d_1 = \frac{\log(S/K) + (r - q + \frac{1}{2}\sigma^2)\tau}{\sigma\sqrt{\tau}}.$$

Training labels use shifted Sobol terminal GBM paths. Validation uses Black-Scholes:

$$C = Se^{-q\tau}\Phi(d_1) - Ke^{-r\tau}\Phi(d_2), \quad d_2 = d_1 - \sigma\sqrt{\tau}.$$

Default proxy: log price fitted by PCHIP in  $d_1$ .

Exact parts: the terminal GBM formula and the Black-Scholes benchmark. Fitted part: the interpolation from finitely many noisy labels. Main failure mode: tiny option values can make raw percentage error unstable, so validation reports both signed price error and relative error with a floor.

### 6.3 Arithmetic Asian option

State:  $S$  and running sum  $A$ . With  $m$  future fixings and total  $N$ ,

$$K_{\text{adj}} = \frac{NK - A - S}{m}.$$

The call payoff can be rewritten as

$$\left( \frac{A + S + \sum_{j=1}^m SG_j}{N} - K \right)_+ = S \frac{m}{N} \left( \frac{1}{m} \sum_{j=1}^m G_j - \frac{K_{\text{adj}}}{S} \right)_+.$$

Thus the tested monthly call can be represented by an adjusted-moneyness coordinate. Training uses Sobol antithetics, likelihood shifts where useful, and a geometric Asian control variate.

Exact parts: the payoff rewrite and deterministic in-the-money or out-of-the-money regions. Unbiased parts: Sobol randomized labels, antithetics, and likelihood weights. Approximate parts: using a one-dimensional adjusted-moneyness fit for the tested family and using MC for the arithmetic-average benchmark.

### 6.4 Ten-asset basket Asian

State: ten spots plus the running basket sum. The proxy starts from a moment-matched lognormal baseline for the final arithmetic basket average, then fits a sparse Chebyshev/PCA correction and a PCHIP residual calibration. PCA summarizes cross-sectional composition without throwing away the level and running-sum information.

Exact parts: first and second basket moments under GBM. Approximate parts: the lognormal moment match and the sparse residual proxy. The moment match is not expected to be exact because an arithmetic basket average is a sum of correlated lognormal variables, and that sum is not lognormal.

### 6.5 American and Bermudan options

At each exercise date  $k$ ,

$$V_k(S) = \max(g_k(S), e^{-r\Delta t} \mathbb{E}[V_{k+1}(S_{k+1}) \mid S_k = S]).$$

The continuation value is estimated by one-step Sobol transitions and fitted with PCHIP. The max with intrinsic value is imposed exactly.

Exact parts: dynamic-programming recursion and the exercise max. Approximate parts: finite exercise grid for Bermudan/American discretization, MC continuation labels, and interpolation of continuation. Benchmarks use a tree or projected finite-difference method to avoid testing the proxy against itself.

## 6.6 Barrier option

For a zero-rebate barrier,

$$V_{\text{in}} + V_{\text{out}} = V_{\text{vanilla}}.$$

Discrete monitoring checks simulated observation dates. Continuous monitoring uses Brownian-bridge survival between observations. For a lower log barrier  $b$ , endpoints  $x, y > b$ , and variance  $\sigma^2 \Delta t$ ,

$$P(\text{survive} \mid x, y) = 1 - \exp\left(-\frac{2(x-b)(y-b)}{\sigma^2 \Delta t}\right).$$

Exact parts: discrete barrier path payoff and in/out parity. Conditional part: the Brownian-bridge survival probability is exact for continuous monitoring under log-GBM between two fixed endpoints. Approximate parts: MC integration and interpolation across spot.

## 6.7 Cliquet and SLV cliquet

Cliquet payoffs are bounded by local and global caps/floors. If accrued value  $a$  and  $m$  remaining coupons imply the final total is certainly below the global floor or above the global cap, the value is exact.

For single-name SLV cliquet, the state is  $(a, S, v)$ . The proxy uses exact tails, bounded-logit targets, and local/sparse Chebyshev features. For three-asset basket cliquet, the fitted proxy is not universally enough for strict max-error control; hard order-statistic cases use a cached Sobol/LR safety proxy.

Exact parts: local coupon clipping, global floor/cap clipping, and tail bounds. Approximate parts: SLV time stepping, sparse fitted surfaces, and any cached MC safety proxy. The safety proxy is slower than a closed-form fitted proxy, but it is still much faster than regenerating a full benchmark for every repeated query.

## 6.8 Autocallable and random payoff

Autocallable valuation is naturally date-indexed. The state is spot at an observation date; the cash-flow rule determines whether the note redeems early or continues. Random payoff tests are synthetic one-dimensional payoffs used to stress interpolation under different shapes.

Exact parts: the call/no-call cash-flow rule at an observation date. Approximate parts: continuation labels and interpolation. The synthetic random payoff tests are not financial products; they are diagnostic shapes used to catch overshoot, poor extrapolation, and kink handling.

## Chapter 7

# Contract Definitions and Payoff Algebra

### 7.1 European call and put

For a call and put with strike  $K$ ,

$$H_{\text{call}} = (S_T - K)_+, \quad H_{\text{put}} = (K - S_T)_+.$$

The proxy problem is one-dimensional because, under constant-parameter GBM, all future uncertainty depends on the current state only through the current spot  $S_t$ . In practice the coordinate  $d_1$ , rather than raw spot, is preferred because equal spacing in  $d_1$  allocates nodes by option delta rather than by currency units of spot.

### 7.2 Arithmetic Asian call

Let  $0 = t_0 < t_1 < \dots < t_{N-1} = T$  be fixing dates. At valuation date  $t_i$ , let

$$A_i = \sum_{j=0}^{i-1} S_{t_j}$$

be the running sum before today's fixing. If today's spot is  $S$ , then the final arithmetic average is

$$\bar{S}_T = \frac{A_i + S + \sum_{k=i+1}^{N-1} S_{t_k}}{N}.$$

The payoff is

$$H = (\bar{S}_T - K)_+.$$

The state is two-dimensional before reduction:  $X_i = (S, A_i)$ . The adjusted-strike algebra in Appendix C explains why the tested call can be fit on a one-dimensional adjusted-moneyness coordinate.

### 7.3 Basket Asian call

For  $d$  assets with basket weights  $w_j$ ,

$$B_{t_k} = \sum_{j=1}^d w_j S_{t_k}^{(j)}, \quad \bar{B}_T = \frac{R_i + B_{t_i} + \sum_{k=i+1}^{N-1} B_{t_k}}{N},$$

where  $R_i$  is the running sum of past basket fixings. The payoff is

$$H = (\bar{B}_T - K)_+.$$

This is no longer reducible to one scalar in general because the future basket distribution depends on the full cross-section of spots and correlations. The code therefore uses moment baselines and PCA features.

### 7.4 Barrier call

For a knock-out call with lower barrier  $L$  and upper barrier  $U$ , the alive indicator is

$$I_T = \mathbf{1}\{L < S_{t_k} < U \text{ for all monitored } t_k\}.$$

The payoff is

$$H_{\text{out}} = I_T(S_T - K)_+.$$

For a knock-in,  $H_{\text{in}} = (1 - I_T)(S_T - K)_+$ , so  $H_{\text{in}} + H_{\text{out}} = H_{\text{vanilla}}$  path by path.

### 7.5 Cliquet payoff

For one underlying, the local return over period  $k$  is

$$R_k = \frac{S_{t_k}}{S_{t_{k-1}}} - 1.$$

The locally clipped coupon is

$$C_k = \text{clip}(R_k, \ell, u).$$

If  $a_i = \sum_{k=1}^i C_k$  is the accrued clipped return at reset  $i$ , the final payoff is

$$H = N_0 \text{ clip}\left(a_i + \sum_{k=i+1}^M C_k, G_L, G_U\right),$$

where  $N_0$  is notional and  $G_L, G_U$  are global floor and cap.

## 7.6 Basket cliquet coupon variants

Let  $R_k^{(j)}$  be the  $j$ -th asset return in period  $k$ . The repository tests several coupon maps. Representative examples are

$$\begin{aligned} C_k^{\text{basket}} &= \text{clip} \left( \sum_j w_j R_k^{(j)}, \ell, u \right), \\ C_k^{\text{worst}} &= \text{clip} \left( \min_j R_k^{(j)}, \ell, u \right), \\ C_k^{\text{best}} &= \text{clip} \left( \max_j R_k^{(j)}, \ell, u \right), \\ C_k^{\text{spread}} &= \text{clip} \left( \sum_j w_j R_k^{(j)}, \ell, u \right) - \lambda_s \text{clip} \left( \max_j R_k^{(j)} - \min_j R_k^{(j)}, s_L, s_U \right) + b \mathbf{1}_{\{\sum_j w_j R_k^{(j)} \geq 0\}}. \end{aligned}$$

Order-statistic coupons are hard because small changes in cross-sectional dispersion can change which asset is best or worst.

## 7.7 Autocallable payoff recursion

At observation date  $k$ , if the note has not called and  $S_{t_k} \geq B_{\text{call}}K$ , the contract redeems. A schematic recursion is

$$V_k(S) = \begin{cases} N_0(1 + c_k), & S \geq B_{\text{call}}K, \\ e^{-r\Delta t} \mathbb{E}[V_{k+1}(S_{k+1}) \mid S_k = S], & S < B_{\text{call}}K. \end{cases}$$

At maturity the redemption rule may include coupon, capital protection, or downside participation.

# Chapter 8

## Implementation Algorithms

### 8.1 Generic proxy construction algorithm

For every instrument, the implementation follows this structure.

1. Define the contract payoff and observation schedule.
2. Prove or choose a Markov state  $X_t$ .
3. Transform the state to dimensionless coordinates: moneyness,  $d_1$ , normalized accrued return, normalized variance, PCA scores, or expected-total cushion.
4. Generate a training state grid. Include broad coverage and extra nodes near exercise, barrier, floor/cap, and low-value transition regions.
5. Generate Monte Carlo labels with Sobol, antithetics, control variates, and likelihood-ratio shifts where they lower variance.
6. Fit a proxy target such as  $\log(V + \varepsilon)$ , a bounded logit, or a baseline residual.
7. Validate on independent states and independent random streams.
8. Report values, signed errors, absolute errors, relative errors with a floor, and benchmark standard errors when available.

### 8.2 European algorithm

1. Fix  $\tau = T - t$ .
2. Choose training nodes  $S_i = S(d_{1,i})$  with  $d_{1,i}$  covering low and high deltas.
3. For each node, estimate  $V_i$  by shifted Sobol terminal paths:

$$S_T^{(i,n)} = S_i \exp\left((r - q - \frac{1}{2}\sigma^2)\tau + \sigma\sqrt{\tau} Z_{i,n}^{\text{shifted}}\right).$$



4. Fit  $y_i = \log(V_i + \varepsilon)$  with PCHIP in  $d_1$ .
5. Predict by  $\hat{V}(S) = \exp(\hat{y}(d_1(S))) - \varepsilon$ , clipped at zero.
6. Benchmark by Black-Scholes.

### 8.3 Asian algorithm

1. For each fixing date, build states  $(S, A)$  in adjusted moneyness and running-average directions.
2. Compute  $K_{\text{adj}} = (NK - A - S)/m$ .
3. Use exact linear/zero regions when  $K_{\text{adj}} \leq 0$  or the option is certainly out of the money.
4. For the active region, simulate future fixings by Sobol antithetic GBM paths.
5. Estimate the arithmetic payoff and geometric payoff on the same paths.
6. Use the exact geometric Asian value as a control variate.
7. Fit normalized value and time value with PCHIP.
8. Validate with an independent high-path benchmark.

### 8.4 Basket Asian algorithm

1. Build states  $X = (S^{(1)}, \dots, S^{(10)}, R)$ .
2. Compute the mean and variance of the final arithmetic basket average.
3. Price the moment-matched lognormal baseline  $B(X)$ .
4. Compute PCA scores of centered log spots to summarize composition.
5. Fit a sparse Chebyshev correction to  $\log((V + \varepsilon)/(B + \varepsilon))$ .
6. Bin residuals against expected-average moneyness and fit a PCHIP residual calibration.
7. Validate against Sobol/LR benchmarks on independent state grids.

### 8.5 American and Bermudan algorithm

1. Set terminal value equal to payoff.
2. Move backward through exercise dates.
3. At each date, simulate one-step transitions from training spots.
4. Evaluate the already-fitted next-date value on simulated next spots.
5. Average and discount to estimate continuation.

6. Fit continuation by PCHIP.
7. Define  $V_k(S) = \max(g_k(S), \hat{C}_k(S))$ .
8. Validate American values against finite differences and Bermudan values against a tree.

## 8.6 Barrier algorithm

1. Build an alive-state spot grid inside the barrier domain.
2. Simulate discrete monitoring paths with Sobol antithetics.
3. For continuous monitoring, multiply each path by Brownian-bridge survival probabilities on each segment.
4. Fit the alive knock-out value by the best validated one-dimensional interpolator.
5. Obtain knock-in values by in/out parity when needed.
6. Validate by independent high-path simulation.

## 8.7 Cliquet and basket cliquet algorithm

1. Represent the state by accrued return and any market variables needed for the next returns.
2. Use exact global floor/cap tails before simulation.
3. Simulate remaining local returns under GBM or SLV.
4. Reuse each simulated future coupon sum for several accrued values when possible.
5. Fit a bounded target in logit space.
6. For high-dimensional order-statistic basket coupons, fall back to cached Sobol/LR safety pricing when fitted surfaces fail max-error targets.

## Chapter 9

# Validation and Timing

### 9.1 Independent benchmark principle

Training labels and benchmark labels use separate random streams. Benchmarks use more paths, closed forms, finite differences, or tree methods when available.

### 9.2 Current reporting

The timing report separates:

- training sample generation;
- proxy fitting after labels are available;
- option-family wall time;
- average wall time per scenario.

The reduced-budget smoke accuracy columns in the timing report are diagnostic only. The product-specific markdown summaries remain the full-budget accuracy source.

## Appendix A

# Line-by-Line Proofs: Conditional Pricing and State Choice

### A.1 Risk-neutral pricing formula

**Theorem A.1** (Risk-neutral valuation, constant short rate). If the discounted asset gains are martingales under  $\mathbb{Q}$  and the payoff  $H$  is integrable, then the time- $t$  model value is

$$V_t = e^{-r(T-t)} \mathbb{E}_{\mathbb{Q}}[H \mid \mathcal{F}_t].$$

**Line 1.** The money-market account is  $B_t = e^{rt}$ .

**Line 2.** Absence of arbitrage in the risk-neutral model says the discounted value process  $B_t^{-1}V_t$  is a  $\mathbb{Q}$ -martingale.

**Line 3.** Therefore

$$B_t^{-1}V_t = \mathbb{E}_{\mathbb{Q}}[B_T^{-1}H \mid \mathcal{F}_t].$$

**Line 4.** Since  $B_t = e^{rt}$  and  $B_T = e^{rT}$ ,

$$e^{-rt}V_t = e^{-rT} \mathbb{E}_{\mathbb{Q}}[H \mid \mathcal{F}_t].$$

**Line 5.** Multiply both sides by  $e^{rt}$  to obtain the formula.

### A.2 When a proxy state is sufficient

**Theorem A.2** (State reduction by conditional distribution). Suppose  $X_t$  is  $\mathcal{F}_t$ -measurable and the conditional distribution of future cash flows  $H$  given  $\mathcal{F}_t$  depends on the past only through  $X_t$ . Then there exists a function  $v$  such that

$$V_t = v(t, X_t).$$

**Line 1.** Define the price by conditional expectation:

$$V_t = e^{-r(T-t)} \mathbb{E}[H \mid \mathcal{F}_t].$$

**Line 2.** By assumption, once  $X_t$  is known, earlier path details do not change the conditional law of  $H$ .

**Line 3.** Therefore  $\mathbb{E}[H \mid \mathcal{F}_t]$  is a function of  $X_t$ .

**Line 4.** Call this function  $g(t, X_t)$ .

**Line 5.** Set  $v(t, X_t) = e^{-r(T-t)}g(t, X_t)$ .

**Line 6.** Hence  $V_t = v(t, X_t)$ .

### A.3 Tower property for backward induction

**Theorem A.3** (Tower property). If  $\mathcal{G} \subseteq \mathcal{F}$ , then

$$\mathbb{E}[\mathbb{E}[Y \mid \mathcal{F}] \mid \mathcal{G}] = \mathbb{E}[Y \mid \mathcal{G}].$$

**Line 1.** The conditional expectation  $\mathbb{E}[Y \mid \mathcal{F}]$  is the best  $\mathcal{F}$ -measurable summary of  $Y$  in the sense of matching integrals over  $\mathcal{F}$ -events.

**Line 2.** Every  $\mathcal{G}$ -event is also an  $\mathcal{F}$ -event because  $\mathcal{G} \subseteq \mathcal{F}$ .

**Line 3.** Therefore replacing  $Y$  by  $\mathbb{E}[Y \mid \mathcal{F}]$  preserves all integrals over  $\mathcal{G}$ -events.

**Line 4.** The conditional expectation with respect to  $\mathcal{G}$  is the unique  $\mathcal{G}$ -measurable variable with those integrals.

**Line 5.** Hence the identity follows.

## Appendix B

# Line-by-Line Proofs: Monte Carlo

### B.1 Unbiased sample mean

**Line 1.** Define  $Y = e^{-r\tau}H$  and  $V = \mathbb{E}_{\mathbb{Q}}[Y]$ .

**Line 2.** Define  $\hat{V}_N = N^{-1} \sum_{n=1}^N Y_n$ .

**Line 3.** Apply linearity:

$$\mathbb{E}[\hat{V}_N] = \mathbb{E}\left[\frac{1}{N} \sum_{n=1}^N Y_n\right] = \frac{1}{N} \sum_{n=1}^N \mathbb{E}[Y_n] = \frac{1}{N} \sum_{n=1}^N V = V.$$

**Line 4.** If paths are independent with variance  $\sigma_Y^2$ ,

$$\text{Var}(\hat{V}_N) = \text{Var}\left(\frac{1}{N} \sum_{n=1}^N Y_n\right) = \frac{1}{N^2} \sum_{n=1}^N \text{Var}(Y_n) = \frac{\sigma_Y^2}{N}.$$

**Line 5.** Replace  $\sigma_Y^2$  by the sample variance to estimate standard error.

### B.2 Likelihood-ratio shift

**Line 1.** Start with  $V = \int h(z)p(z) dz$ .

**Line 2.** Choose a proposal  $q$  with  $q(z) > 0$  wherever  $h(z)p(z) \neq 0$ .

**Line 3.** Multiply and divide by  $q$ :

$$V = \int h(z) \frac{p(z)}{q(z)} q(z) dz = \mathbb{E}_q[h(Z)L(Z)], \quad L(Z) = \frac{p(Z)}{q(Z)}.$$

**Line 4.** For  $p(z) = \phi(z)$ ,  $q(z) = \phi(z - \theta)$ ,

$$L(z) = \frac{\exp(-z^2/2)}{\exp(-(z - \theta)^2/2)} = \exp\left(-\theta z + \frac{\theta^2}{2}\right).$$

**Line 5.** For independent shocks, likelihood ratios multiply:

$$L(z_1, \dots, z_m) = \prod_{j=1}^m \exp\left(-\theta_j z_j + \frac{\theta_j^2}{2}\right).$$

### B.3 Mixture likelihood ratio

**Line 1.** Let the original density be  $p(z)$ .

**Line 2.** Choose shifted proposal densities  $q_m(z)$  and probabilities  $\pi_m$  with  $\sum_m \pi_m = 1$ .

**Line 3.** Define the mixture density

$$q(z) = \sum_{m=1}^M \pi_m q_m(z).$$

**Line 4.** Require  $q(z) > 0$  wherever  $h(z)p(z) \neq 0$ .

**Line 5.** Multiply and divide by  $q$ :

$$\int h(z)p(z) \, dz = \int h(z) \frac{p(z)}{q(z)} q(z) \, dz.$$

**Line 6.** Therefore, if  $Z \sim q$ ,

$$\mathbb{E}_p[h(Z)] = \mathbb{E}_q[h(Z)L(Z)], \quad L(Z) = \frac{p(Z)}{q(Z)}.$$

**Line 7.** To simulate from  $q$ , first draw a component  $M$  with probability  $\pi_m$ , then draw  $Z \sim q_M$ .

**Line 8.** The likelihood ratio must use the full mixture denominator  $q(z)$ , not only the chosen component density  $q_M(z)$ , because the sampled value  $z$  could also have been generated by other components.

### B.4 Antithetic estimator

**Line 1.** Standard normal symmetry gives  $Z \stackrel{d}{=} -Z$ .

**Line 2.** Define  $A = (f(Z) + f(-Z))/2$ .

**Line 3.** Then

$$\mathbb{E}[A] = \frac{\mathbb{E}[f(Z)] + \mathbb{E}[f(-Z)]}{2} = \frac{\mathbb{E}[f(Z)] + \mathbb{E}[f(Z)]}{2} = \mathbb{E}[f(Z)].$$

**Line 4.** The variance is

$$\text{Var}(A) = \frac{1}{2} \text{Var}(f(Z)) + \frac{1}{2} \text{Cov}(f(Z), f(-Z)).$$

**Line 5.** The pairing helps whenever the covariance term is negative.

## Appendix C

# Line-by-Line Proofs: GBM and European Options

### C.1 GBM solution

**Line 1.** Start from  $dS_t/S_t = (r - q) dt + \sigma dW_t$ .

**Line 2.** Apply Ito's lemma to  $f(S) = \log S$ :

$$d \log S_t = \frac{1}{S_t} dS_t - \frac{1}{2} \frac{1}{S_t^2} (dS_t)^2.$$

**Line 3.** Since  $(dS_t)^2 = \sigma^2 S_t^2 dt$ ,

$$d \log S_t = (r - q - \frac{1}{2} \sigma^2) dt + \sigma dW_t.$$

**Line 4.** Integrate over  $[t, T]$ :

$$\log(S_T/S_t) = (r - q - \frac{1}{2} \sigma^2) \tau + \sigma(W_T - W_t).$$

**Line 5.** Use  $W_T - W_t = \sqrt{\tau}Z$ ,  $Z \sim N(0, 1)$ , and exponentiate.

### C.2 Black-Scholes call

**Line 1.** Risk-neutral pricing:

$$C = e^{-r\tau} \mathbb{E}[(S_T - K)_+].$$

**Line 2.** Split the positive part:

$$\mathbb{E}[(S_T - K)_+] = \mathbb{E}[S_T \mathbf{1}_{\{S_T > K\}}] - K \mathbb{Q}(S_T > K).$$

**Line 3.** The event  $S_T > K$  is  $Z > -d_2$ , so  $\mathbb{Q}(S_T > K) = \Phi(d_2)$ .

**Line 4.** Completing the square gives

$$\mathbb{E}[S_T \mathbf{1}_{\{S_T > K\}}] = S e^{(r-q)\tau} \Phi(d_1).$$

**Line 5.** Substitute and discount:

$$C = S e^{-q\tau} \Phi(d_1) - K e^{-r\tau} \Phi(d_2).$$



## Appendix D

# Line-by-Line Proofs: Asian State Reduction

### D.1 Adjusted moneyness

**Line 1.** At a valuation date, write future fixings as  $SG_1, \dots, SG_m$ .

**Line 2.** The arithmetic average is

$$\bar{S} = \frac{A + S + S \sum_{j=1}^m G_j}{N}.$$

**Line 3.** Define

$$K_{\text{adj}} = \frac{NK - A - S}{m}.$$

**Line 4.** Rewrite the call payoff:

$$(\bar{S} - K)_+ = \frac{1}{N} \left( A + S + S \sum_{j=1}^m G_j - NK \right)_+ = \frac{m}{N} \left( S \frac{1}{m} \sum_{j=1}^m G_j - K_{\text{adj}} \right)_+.$$

**Line 5.** Factor out  $S > 0$ :

$$(\bar{S} - K)_+ = S \frac{m}{N} \left( \frac{1}{m} \sum_{j=1}^m G_j - \frac{K_{\text{adj}}}{S} \right)_+.$$

**Line 6.** Under GBM the growth factors  $G_j$  do not depend on the current spot level  $S$ .

**Line 7.** Therefore  $V(S, A) = SF(K_{\text{adj}}/S, t)$ , which justifies a one-dimensional adjusted-moneyness proxy for the tested call setup.

## D.2 Geometric Asian control variate

For future fixing times  $t_1, \dots, t_m$ , define the geometric average

$$G = \left( \prod_{j=1}^m S_{t_j} \right)^{1/m}.$$

**Line 1.** Take logs:

$$\log G = \frac{1}{m} \sum_{j=1}^m \log S_{t_j}.$$

**Line 2.** Under GBM, each  $\log S_{t_j}$  is normal.

**Line 3.** A linear combination of jointly normal random variables is normal.

**Line 4.** Therefore  $\log G \sim N(\mu_G, \sigma_G^2)$ .

**Line 5.** The mean is

$$\mu_G = \log S_0 + \frac{1}{m} \sum_{j=1}^m (r - q - \tfrac{1}{2}\sigma^2)t_j.$$

**Line 6.** The variance is

$$\sigma_G^2 = \frac{\sigma^2}{m^2} \sum_{i=1}^m \sum_{j=1}^m \min(t_i, t_j),$$

because  $\text{Cov}(W_{t_i}, W_{t_j}) = \min(t_i, t_j)$ .

**Line 7.** Since  $G$  is lognormal, the call value  $\mathbb{E}[(G - K)_+]$  has the same algebraic form as Black-Scholes with log-mean  $\mu_G$  and log-variance  $\sigma_G^2$ .

**Line 8.** This exact geometric Asian value can be used as the known mean  $m_C$  in the control-variate formula.

# Appendix E

## Line-by-Line Proofs: Barriers

### E.1 In/out parity

**Line 1.** Split all paths into two disjoint events: hit and no hit.

**Line 2.** On hit paths, the knock-in pays the vanilla payoff and the knock-out pays zero.

**Line 3.** On no-hit paths, the knock-out pays the vanilla payoff and the knock-in pays zero.

**Line 4.** Therefore, path by path,

$$H_{\text{in}} + H_{\text{out}} = H_{\text{vanilla}}.$$

**Line 5.** Discount and take expectations:

$$V_{\text{in}} + V_{\text{out}} = V_{\text{vanilla}}.$$

### E.2 Brownian-bridge lower-barrier survival

**Theorem E.1** (Brownian-bridge crossing probability). Let  $X_s$  be a Brownian bridge from  $x$  to  $y$  over length  $\Delta t$ , with diffusion variance  $\sigma^2$ . If  $x > b$  and  $y > b$ , then

$$\mathbb{Q}\left(\min_{0 \leq s \leq \Delta t} X_s \leq b \mid X_0 = x, X_{\Delta t} = y\right) = \exp\left(-\frac{2(x-b)(y-b)}{\sigma^2 \Delta t}\right).$$

**Line 1.** Condition on log endpoints  $X_0 = x$ ,  $X_{\Delta t} = y$ , both above lower barrier  $b$ .

**Line 2.** By the reflection principle, the conditional probability of crossing below  $b$  is

$$\exp\left(-\frac{2(x-b)(y-b)}{\sigma^2 \Delta t}\right).$$

**Line 3.** Survival is one minus crossing:

$$P_{\text{surv}} = 1 - \exp\left(-\frac{2(x-b)(y-b)}{\sigma^2 \Delta t}\right).$$

**Line 4.** Conditional survival over many intervals is the product of interval survival probabilities.

## Appendix F

# Line-by-Line Proofs: Exercise and Cliquet Bounds

### F.1 American recursion

**Line 1.** At final date  $M$ ,  $V_M(S) = g_M(S)$ .

**Line 2.** Suppose  $V_{k+1}$  is known.

**Line 3.** Continuing from date  $k$  has value

$$C_k(S) = e^{-r\Delta t} \mathbb{E}[V_{k+1}(S_{k+1}) \mid S_k = S].$$

**Line 4.** Stopping immediately has value  $g_k(S)$ .

**Line 5.** The holder chooses the larger:

$$V_k(S) = \max(g_k(S), C_k(S)).$$

**Line 6.** Backward induction proves optimality at every date.

### F.2 Cliquet exact tails

**Line 1.** Let current accrued return be  $a$ , with  $m$  future local coupons.

**Line 2.** Each future coupon lies in  $[\ell, u]$ .

**Line 3.** Therefore final accumulated return lies in

$$a + \sum_{j=1}^m c_j \in [a + m\ell, a + mu].$$

**Line 4.** If  $a + mu$  is below the global floor, the payoff is exactly the global floor.

**Line 5.** If  $a + m\ell$  is above the global cap, the payoff is exactly the global cap.

**Line 6.** These states need no regression and no simulation.

# Appendix G

## Line-by-Line Proofs: PCHIP

### G.1 Unique Hermite cubic

**Line 1.** A cubic has four coefficients.

**Line 2.** Endpoint value constraints give two linear equations.

**Line 3.** Endpoint slope constraints give two more linear equations.

**Line 4.** The Hermite basis constructs one solution:

$$H(t) = y_i h_{00}(t) + h_i d_i h_{10}(t) + y_{i+1} h_{01}(t) + h_i d_{i+1} h_{11}(t).$$

**Line 5.** If two cubics satisfied all four conditions, their difference would have double roots at both endpoints.

**Line 6.** A nonzero cubic cannot have four roots counting multiplicity, so the solution is unique.

### G.2 Why PCHIP avoids artificial extrema

**Theorem G.1** (Fritsch-Carlson monotonicity criterion). For a cubic Hermite interpolant on a monotone interval, if the normalized endpoint slopes satisfy the Fritsch-Carlson inequalities, then the derivative does not change sign inside the interval. Hence the cubic segment is monotone on that interval.

**Line 1.** Let  $\delta_i = (y_{i+1} - y_i)/h_i$ .

**Line 2.** If neighboring secants change sign, PCHIP sets the node slope to zero.

**Line 3.** If neighboring secants have the same sign, PCHIP uses the weighted harmonic mean

$$d_i = \frac{w_1 + w_2}{w_1/\delta_{i-1} + w_2/\delta_i}.$$

**Line 4.** This slope has the same sign as the neighboring secants and lies between their magnitudes.

- Line 5.** On a monotone interval, the derivative of the Hermite cubic is a quadratic whose normalized endpoint slopes satisfy the Fritsch-Carlson monotonicity inequalities.
- Line 6.** Therefore the derivative does not change sign inside the interval.
- Line 7.** Hence the interpolant cannot create a new local maximum or minimum between two monotone data points.

## Appendix H

# Line-by-Line Proofs: Sparse Chebyshev Ridge

### H.1 Chebyshev boundedness

**Line 1.** Every  $x \in [-1, 1]$  can be written as  $x = \cos \theta$ .

**Line 2.** Define  $T_n(x) = \cos(n\theta)$ .

**Line 3.** Since  $|\cos(n\theta)| \leq 1$ ,

$$|T_n(x)| \leq 1, \quad x \in [-1, 1].$$

**Line 4.** Scaling state variables to  $[-1, 1]$  keeps basis columns numerically bounded.

### H.2 Ridge normal equation

**Line 1.** Objective:

$$J(\beta) = \|A\beta - y\|_2^2 + \lambda \|\Gamma\beta\|_2^2.$$

**Line 2.** Differentiate:

$$\nabla J(\beta) = 2A^\top(A\beta - y) + 2\lambda\Gamma^\top\Gamma\beta.$$

**Line 3.** Set the gradient to zero:

$$A^\top A\beta - A^\top y + \lambda\Gamma^\top\Gamma\beta = 0.$$

**Line 4.** Rearrange:

$$(A^\top A + \lambda\Gamma^\top\Gamma)\beta = A^\top y.$$

**Line 5.** The left matrix is positive semidefinite, and ridge makes unsupported directions stable.

# Appendix I

## Line-by-Line Proofs: Bounded Targets and Metrics

### I.1 Logit bounded target

**Line 1.** Suppose the price must lie in  $[L, U]$ .

**Line 2.** Fit an unconstrained number  $y$  and map it through the logistic function:

$$u(y) = \frac{1}{1 + e^{-y}}.$$

**Line 3.** Because  $e^{-y} > 0$ ,  $0 < u(y) < 1$ .

**Line 4.** Rescale:

$$\hat{V} = L + (U - L)u(y).$$

**Line 5.** Therefore  $L < \hat{V} < U$ .

### I.2 Relative-error floor

**Line 1.** Define

$$e_{\text{rel}} = \frac{|\hat{V} - V|}{\max(|V|, \eta)}, \quad \eta > 0.$$

**Line 2.** If  $|V| \geq \eta$ , this is ordinary relative error.

**Line 3.** If  $|V| < \eta$ , then

$$e_{\text{rel}} = \frac{|\hat{V} - V|}{\eta}.$$

**Line 4.** Thus a nearly zero benchmark cannot create an infinite percentage error.



## Appendix J

# Line-by-Line Proofs: Basket Moments and PCA

### J.1 Basket Asian first moment

Let the basket fixing at time  $t_k$  be

$$B_{t_k} = \sum_{j=1}^d w_j S_{t_k}^{(j)}.$$

Under correlated risk-neutral GBM,

$$\mathbb{E}[S_{t_k}^{(j)} \mid S_0^{(j)}] = S_0^{(j)} e^{(r-q_j)t_k}.$$

Therefore

$$\mathbb{E}[B_{t_k}] = \sum_{j=1}^d w_j S_0^{(j)} e^{(r-q_j)t_k}.$$

For the final arithmetic basket average

$$\bar{B} = \frac{R_i + B_{t_i} + \sum_{k=i+1}^{N-1} B_{t_k}}{N},$$

linearity gives

$$\mathbb{E}[\bar{B}] = \frac{R_i + B_{t_i} + \sum_{k=i+1}^{N-1} \mathbb{E}[B_{t_k}]}{N}.$$

### J.2 Basket Asian second moment

For two assets  $j, \ell$  and times  $t_a, t_b$ , correlated GBM gives

$$\mathbb{E}[S_{t_a}^{(j)} S_{t_b}^{(\ell)}] = S_0^{(j)} S_0^{(\ell)} e^{(r-q_j)t_a + (r-q_\ell)t_b + \rho_{j\ell} \sigma_j \sigma_\ell \min(t_a, t_b)}.$$

Then

$$\mathbb{E}[B_{t_a} B_{t_b}] = \sum_{j=1}^d \sum_{\ell=1}^d w_j w_\ell \mathbb{E}[S_{t_a}^{(j)} S_{t_b}^{(\ell)}].$$

Finally,

$$\text{Var}(\bar{B}) = \mathbb{E}[\bar{B}^2] - \mathbb{E}[\bar{B}]^2,$$

where  $\mathbb{E}[\bar{B}^2]$  is obtained by summing  $\mathbb{E}[B_{t_a} B_{t_b}]$  over all future fixing pairs and adding deterministic running-sum terms.

### J.3 Moment-matched lognormal baseline

**Line 1.** Suppose a positive average has mean  $m$  and variance  $v$ .

**Line 2.** Approximate it by  $L = e^Y$ ,  $Y \sim N(\mu_L, s_L^2)$ .

**Line 3.** The lognormal mean and variance are

$$\mathbb{E}[L] = e^{\mu_L + s_L^2/2}, \quad \text{Var}(L) = e^{2\mu_L + s_L^2} (e^{s_L^2} - 1).$$

**Line 4.** Set  $e^{\mu_L + s_L^2/2} = m$ .

**Line 5.** Divide the variance equation by  $m^2$ :

$$\frac{v}{m^2} = e^{s_L^2} - 1.$$

**Line 6.** Hence

$$s_L^2 = \log(1 + v/m^2), \quad \mu_L = \log m - \frac{1}{2}s_L^2.$$

**Line 7.** Use the Black-Scholes-style truncated-lognormal formula to value  $(L - K)_+$ .

### J.4 PCA optimality

**Theorem J.1** (Rayleigh-Ritz/PCA optimality). For a centered random vector with covariance matrix  $\Sigma$ , the rank- $k$  orthogonal projection minimizing mean-square reconstruction error projects onto the eigenvectors of  $\Sigma$  with the  $k$  largest eigenvalues.

**Line 1.** Let  $R$  be a centered random vector with covariance  $\Sigma$ .

**Line 2.** A rank- $k$  orthogonal projection  $P$  has reconstruction error

$$\mathbb{E}[\|R - PR\|^2] = \text{tr}(\Sigma) - \text{tr}(P\Sigma).$$

**Line 3.** Minimizing error is therefore equivalent to maximizing  $\text{tr}(P\Sigma)$ .

**Line 4.** The Rayleigh-Ritz theorem says this trace is maximized by projecting onto the eigenvectors with the  $k$  largest eigenvalues.

**Line 5.** Therefore PCA gives the best  $k$ -dimensional linear summary in mean-square reconstruction error.

## Appendix K

# Line-by-Line Proofs: Sobol, Design, and Validation

### K.1 Why powers of two are used

Sobol sequences are digital nets in base two. The first  $2^m$  points have the balance property for elementary binary boxes. If we stop at an arbitrary  $N$ , the last partial block may break this balance. The implementation therefore rounds a target path count to

$$N_{\text{Sobol}} = 2^{\lceil \log_2 N_{\text{target}} \rceil}.$$

This is not an unbiasedness theorem. It is a numerical integration design rule.

### K.2 Common random numbers

**Line 1.** At state  $x_i$ , the label is  $V(x_i) = \mathbb{E}[h(x_i, Z)]$ .

**Line 2.** Using the same  $Z_n$  for several states gives

$$\hat{V}(x_i) = \frac{1}{N} \sum_{n=1}^N h(x_i, Z_n).$$

**Line 3.** For each fixed  $i$ , the marginal distribution of  $Z_n$  is unchanged.

**Line 4.** Therefore each  $\hat{V}(x_i)$  remains unbiased.

**Line 5.** The errors across nearby states become correlated, which often makes label noise smoother and easier for interpolation to absorb.

### K.3 Benchmark independence

If the same paths are used for training labels and benchmark labels, the measured proxy error can be artificially small because both contain the same random noise. Independent benchmark

streams estimate

$$\widehat{V}_{\text{proxy}}(x) - V_{\text{model}}(x)$$

with less leakage. When a closed-form, finite-difference, or tree benchmark exists, it is preferred because it removes benchmark Monte Carlo noise.

## Appendix L

# Line-by-Line Proofs: SLV and Basket Cliquet

### L.1 Variance process discretization

The SLV cliquet code uses an Euler-type variance update with positivity protection. A schematic Heston-style variance step is

$$v_{t+\Delta t} = \max\left(v_t + \kappa(\theta - v_t)\Delta t + \xi\sqrt{v_t^+}\sqrt{\Delta t}Z_v, 0\right).$$

The max is not an exact theorem; it is a numerical safeguard. The consequence is that benchmark quality must be checked by increasing time steps or comparing against a stronger scheme when precision matters.

### L.2 Local volatility multiplier

A local volatility multiplier  $L(S)$  changes the spot diffusion to

$$\frac{dS_t}{S_t} = (r - q)dt + L(S_t)\sqrt{v_t}dW_t.$$

The pricing theorem still applies once the simulator defines the risk-neutral dynamics. The proxy must include enough state information, at least  $S_t$  and  $v_t$ , because the future law depends on both.

### L.3 Grouped accrued labels

For cliquets, suppose a simulated future path gives a remaining coupon sum  $C^{(n)}$ . For accrued values  $a_1, \dots, a_L$ ,

$$H_\ell^{(n)} = N_0 \text{clip}(a_\ell + C^{(n)}, G_L, G_U).$$

One simulation of  $C^{(n)}$  prices all  $L$  accrued layers. This is exact path reuse, not approximation, because the same future path sum is valid for every starting accrued value.

## L.4 Why basket order statistics are hard

For a worst-of coupon,

$$C_k = \text{clip}(\min_j R_k^{(j)}, \ell, u).$$

The identity of the minimum can change abruptly as the state changes. This creates nonsmooth interaction surfaces in spot and variance. A low-degree global polynomial can miss these switching surfaces. That is why the methodology allows a slower cached Sobol/LR safety proxy for strict max-error cases.

# References

- Black, F. and Scholes, M. (1973). The Pricing of Options and Corporate Liabilities. *Journal of Political Economy*.
- Fritsch, F. N. and Carlson, R. E. (1980). Monotone Piecewise Cubic Interpolation. *SIAM Journal on Numerical Analysis*.
- Fritsch, F. N. and Butland, J. (1984). A Method for Constructing Local Monotone Piecewise Cubic Interpolants. *SIAM Journal on Scientific and Statistical Computing*.
- Hoerl, A. E. and Kennard, R. W. (1970). Ridge Regression: Biased Estimation for Nonorthogonal Problems. *Technometrics*.
- Glasserman, P. (2004). *Monte Carlo Methods in Financial Engineering*. Springer.
- Glasserman, P. and Staum, J. (2001). Conditioning on One-Step Survival for Barrier Option Simulations. *Operations Research*.
- Broadie, M., Glasserman, P., and Kou, S. G. (1997). A Continuity Correction for Discrete Barrier Options. *Mathematical Finance*.